

Linux/Android

内核 Heap Spray的几种姿势

盘古实验室 River



PANGU TEAM



DEFCON
GROUP 0531
HACKER COMMUNITY



PANGU TEAM



内容概要

- For newbies
 - 基础简介, 环境搭建
 - SLUB细节和特性
- 讨论内核Heap Spray
 - 两个实例
 - 几种spray的实际应用



DEFCON
GROUP 0531
HACKER COMMUNITY




For newbies

Android/Linux 内核提权

- 获取Root权限（通过漏洞）

- \$ -> #
- 改变uid, gid

```
river@ubuntu:~$ id  
uid=1000(river) gid=1000(river) groups=1000
```



```
root@ubuntu:~# id  
uid=0(root) gid=0(root) groups=0(root)
```

Linux



- 开源
 - 仅仅是内核
 - 发行版(Ubuntu, Arch, Debian,CentOS...) != Linux
 - GNU+Linux = GNU/Linux --> 发行版(Ubuntu, Arch, Debian,CentOS...)
- www.kernel.org

The Linux Kernel Archives



[About](#) [Contact us](#) [FAQ](#) [Releases](#) [Signatures](#) [Site news](#)

Protocol	Location
HTTP	https://www.kernel.org/pub/
GIT	https://git.kernel.org/
RSYNC	rsync://rsync.kernel.org/pub/

Latest Stable Kernel:



4.17.2

mainline:	4.18-rc1	2018-06-16	[tarball]	[patch]	[view diff]	[browse]	
stable:	4.17.2	2018-06-16	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]
stable:	4.16.17	2018-06-20	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]
longterm:	4.14.51	2018-06-20	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]
longterm:	4.9.109	2018-06-16	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]
longterm:	4.4.138	2018-06-16	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]
longterm:	4.1.52 [EOL]	2018-05-28	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]
longterm:	3.18.113 [EOL]	2018-06-13	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]
longterm:	3.16.57	2018-06-16	[tarball]	[pgp] [patch] [inc. patch]	[view diff]	[browse]	[changelog]

Android



Android

- 开源

- 内核采用Linux(ion,ashmem...), 外加SoC厂商代码
- AOSP+(Linux) = Android
- 高通, 海思, 三星, 博通, 英伟达, mtk, mstar

- <https://android.googlesource.com>

- androidxref.com




DEFCON
GROUP 0531
HACKER COMMUNITY



内核

- 内核管理硬件资源
 - CPU资源、存储资源、I/O资源等
- 内核为用户态提供服务: `syscall()`
 - fs
 - net
 - ptrace
 -



syscall

pipe()

rename()

ffffffffffffbd261630	T	sys_pipe2
ffffffffffffbd2616e0	T	sys_pipe
ffffffffffffbd2691f0	T	sys_mknodat
ffffffffffffbd269400	T	sys_mknod
ffffffffffffbd269610	T	sys_mkdirat
ffffffffffffbd269710	T	sys_mkdir
ffffffffffffbd269810	T	sys_rmdir
ffffffffffffbd269830	T	sys_unlinkat
ffffffffffffbd269860	T	sys_unlink
ffffffffffffbd269880	T	sys_symlinkat
ffffffffffffbd269990	T	sys_symlink
ffffffffffffbd269a90	T	sys_linkat
ffffffffffffbd269d10	T	sys_link
ffffffffffffbd269f30	T	sys_renameat2
ffffffffffffbd26a500	T	sys_renameat
ffffffffffffbd26a8f0	T	sys_rename
ffffffffffffbd26bb30	T	sys_fcntl



环境搭建



- git clone
 - <https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux-stable.git>
 - git tags、checkout 、 git stash
- 编译

环境搭建



- qemu运行内核

```
#!/bin/bash
qemu-system-x86_64 -hda $IMG -m 1024M -net nic -net
:1314-:22 --enable-kvm -kernel $KERNEL -append "root
panic panic=1 quiet" -smp cores=4,threads=2,sockets=
ic -cpu kvm64
```

环境搭建



Android

- git clone
 - msm 、 华为、 mtk、 三星
 - goldfish !
- 编译
 - ndk toolchain、 aarch64

环境搭建



Android

- emulator
 - aosp里的build-kernel.sh 脚本（可编译适配的内核）
- 实体机
 - msm ...
 - abooting 替换内核

调试



- gdb
 - i r、 x、 b、 disas、 watch、 s/n、 si/ni
- kasan、 slub_info
- 插log
 - WARN_ON (1) 、 printk

调试



- strace
- ftrace
- 串口


内核漏洞类型

- oob
- stack/heap overflow
- use-after-free
- double free
- race condition




overflow

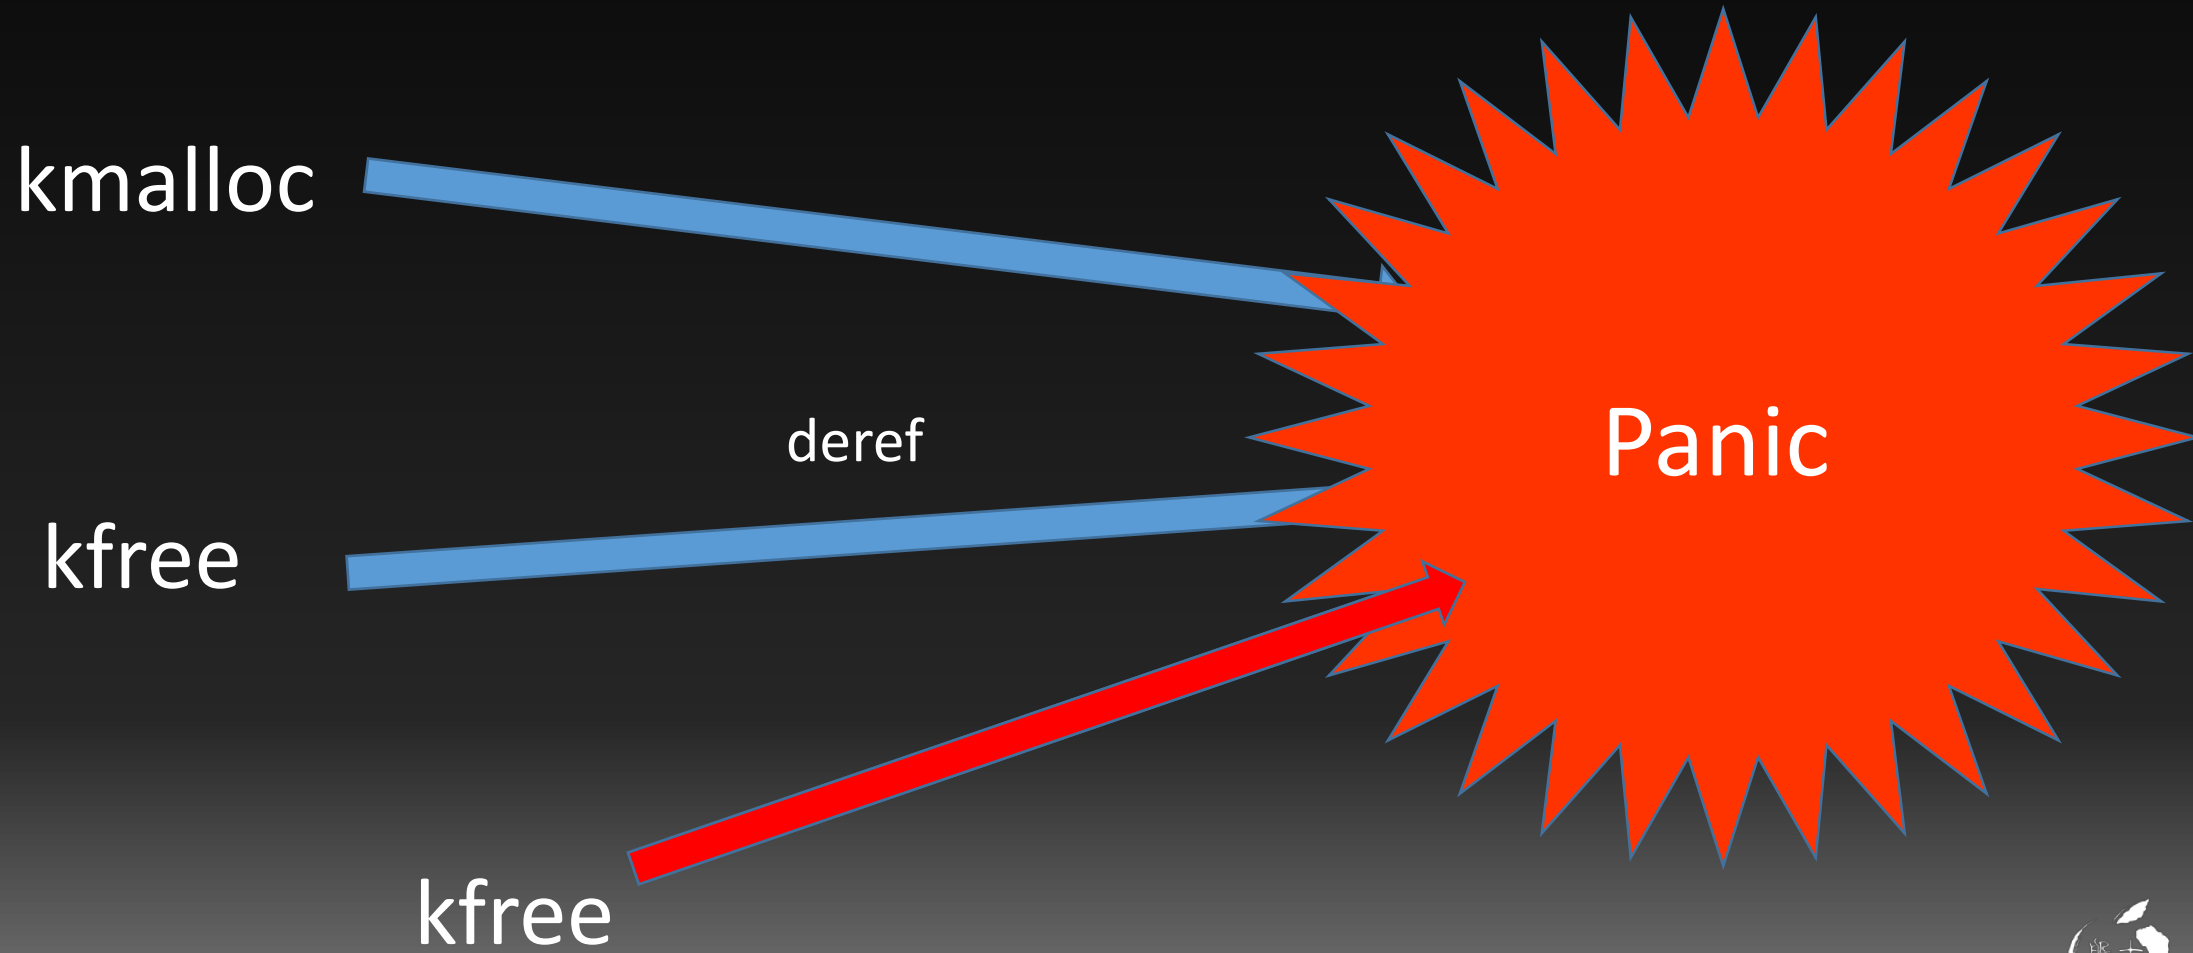
buffer



UAF



Double Free



SLUB

- 由buddy system 提供页支持

- slob、slab、slub



- slub 里划分为多个obj空间

- 专为小内存分配而生

- kmem_cache_create

- kmalloc-64, kmalloc-128, kmalloc-256



```

62 struct kmem_cache {
63     struct kmem_cache_cpu __percpu *cpu_slab;
64     /* Used for retrieving partial slabs etc */
65     unsigned long flags;
66     unsigned long min_partial;
67     int size; /* The size of an object including meta data */
68     int object_size; /* The size of an object without meta data */
69     int offset; /* Free pointer offset. */
70     int cpu_partial; /* Number of per cpu partial objects to keep around */
71     struct kmem_cache_order_objects oo;
72
73     /* Allocation and freeing of slabs */
74     struct kmem_cache_order_objects max;
75     struct kmem_cache_order_objects min;
76     gfp_t allocflags; /* gfp flags to use on each alloc */
77     int refcount; /* Refcount for slab cache destroy */
78     void (*ctor)(void *);
79     int inuse; /* Offset to metadata */
80     int align; /* Alignment */
81     int reserved; /* Reserved bytes at the end of slabs */
82     const char *name; /* Name (only for display!) */
83     struct list_head list; /* List of slab caches */
84 #ifdef CONFIG_SYSFS
85     struct kobject kobj; /* For sysfs */
86 #endif
87 #ifdef CONFIG_MEMCG_KMEM
88     struct memcg_cache_params *memcg_params;
89     int max_attr_size; /* for propagation, maximum size of a stored attr */
90 #ifdef CONFIG_SYSFS
91     struct kset *memcg_kset;
92 #endif
93 #endif
94
95 #ifdef CONFIG_NUMA
96     /*
97      * Defragmentation by allocating from a remote node.
98      */
99     int remote_node_defrag_ratio;
100 #endif
101     struct kmem_cache_node *node[MAX_NUMNODES];
102 };

```

一个per cpu变量，对于每个cpu来说，相当于一个本地内存缓存池。当分配内存的时候优先从本地cpu分配内存以保证cache的命中率



freelist

```
40 struct kmem_cache_cpu {
41     void **freelist;    /* Pointer to next available object */
42     unsigned long tid;  /* Globally unique transaction id */
43     struct page *page;  /* The slab from which we are allocating */
44     struct page *partial; /* Partially allocated frozen slabs */
45 #ifdef CONFIG_SLUB_STATS
46     unsigned stat[NR_SLUB_STAT_ITEMS];
47 #endif
48 };
```



SLUB

- 故而，slab内存块的释放和CPU id息息相关
 - refill或者shape heap时，exp的某些线程注意绑定某CPU id

```
128 void bind_cpu(int cpuid) {  
129     int ret;  
130     cpu_set_t mask;  
131     CPU_ZERO(&mask);  
132     CPU_SET(cpuid, &mask);  
133     ret = sched_setaffinity(0, sizeof(mask), &mask);  
134     CHKERR(ret, "sched_setaffinity set failed ...");  
135 }  
136
```

缓解措施

- Linux
 - smep、smap、no namespace、kaslr
- Android
 - selinux、pxn、pan、kaslr

内核 Heap Spray

Heap Spray

- uaf/double-free、 race
 - 有时利用需要refill obj
- 堆溢出/泄露类
 - shape heap/fengshui


UAF / Double Free

- uaf/double free
 - CVE-2015-3636
 - CVE-2017-0403
 - CVE-2016-6787
 - CVE-2017-8824
 - CVE-2017-8890
 - CVE-2017-17053
 - CVE-2017-6074


Race Condition

- race condition
 - CVE-2016-1805
 - CVE-2017-10661
 - CVE-2017-7533 (heap overflow or UAF)
 -

数据流劫持(refill)



数据流劫持(shape heap)



实例分析1

- Double free: CVE-2017-8890
 - 指针残留导致的UAF

```
1 diff --git a/net/ipv4/inet_connection_sock.c b/net/ipv4/inet_connection_sock.c
2 index 5e313c1..1054d33 100644
3 --- a/net/ipv4/inet_connection_sock.c
4 +++ b/net/ipv4/inet_connection_sock.c
5 @@ -794,6 +794,8 @@ struct sock *inet_csk_clone_lock(const struct sock *sk,
6      /* listeners have SOCK_RCU_FREE, not the children */
7      sock_reset_flag(newsk, SOCK_RCU_FREE);
8
9 +    inet_sk(newsk)->mc_list = NULL;
10 +
11     newsk->sk_mark = inet_rsk(req)->ir_mark;
12     atomic64_set(&newsk->sk_cookie,
13         atomic64_read(&inet_rsk(req)->ir_cookie));
```

实例分析1

- `accept4()` -> -> `inet_csk_clone_lock()` -> `sock_clone_lock`
- 产生一个新的sock

实例分析1

```
672 struct sock *inet_csk_clone_lock(const struct sock *sk,  
673                                 const struct request_sock *req,  
674                                 const gfp_t priority)  
675 {  
676     struct sock *newsk = sk_clone_lock(sk, priority);  
677  
678     if (newsk != NULL) {  
679         struct inet_connection_sock *newicsk = inet_csk(newsk);  
680  
681         newsk->sk_state = TCP_SYN_RECV;  
682         newicsk->icsk_bind_hash = NULL;  
683  
684         inet_sk(newsk)->inet_dport = inet_rsk(req)->ir_rmt_port;  
685         inet_sk(newsk)->inet_num = inet_rsk(req)->ir_num;  
686         inet_sk(newsk)->inet_sport = htons(inet_rsk(req)->ir_num);  
687         newsk->sk_write_space = sk_stream_write_space;  
688  
689         newsk->sk_mark = inet_rsk(req)->ir_mark;  
690  
691         newicsk->icsk_retransmits = 0;  
692         newicsk->icsk_backoff = 0;  
693         newicsk->icsk_probes_out = 0;  
694  
695         /* Deinitialize accept_queue to trap illegal accesses. */  
696         memset(&newicsk->icsk_accept_queue, 0, sizeof(newicsk->icsk_accept_queue));  
697  
698         security_inet_csk_clone(newsk, req);  
699     }  
700     return newsk;  
701 }
```



实例分析1

```
1475 struct sock *sk_clone_lock(const struct sock *sk, const gfp_t priority)
1476 {
1477     struct sock newsk;
1478     bool is_charged = true;
1479
1480     newsk = sk_prot_alloc(sk->sk_prot, priority, sk->sk_family);
1481     if (newsk != NULL) {
1482         struct sk_filter *filter;
1483
1484         sock_copy(newsk, sk);
1485
1486         /* SANITY */
1487         get_net(sock_net(newsk));
1488         sk_node_init(&newsk->sk_node);
1489         sock_lock_init(newsk);
1490         bh_lock_sock(newsk);
1491         newsk->sk_backlog.head = newsk->sk_backlog.tail = NULL;
1492         newsk->sk_backlog.len = 0;
```



实例分析1

- mc_list 指向同一块内存

```
154 struct inet_sock {
155     /* sk and pinet6 has to be the first two members of inet_sock */
156     struct sock      sk;
157     #if IS_ENABLED(CONFIG_IPV6)
158     struct ipv6_pinfo *pinet6;
159     #endif
160     /* Socket demultiplex comparisons on incoming packets. */
161     #define inet_daddr      sk->__sk_common.skc_daddr
162     #define inet_rcv_saddr  sk->__sk_common.skc_rcv_saddr
163     #define inet_dport      sk->__sk_common.skc_dport
164     #define inet_num        sk->__sk_common.skc_num
165
166     // ---- snip ----
167
168     be32      mc_addr;
169     struct ip_mc_socklist __rcu *mc_list;
170     struct inet_cork_full cork;
171 };
172
```

实例分析1

- mc_list 指向同一块内存，但copy时候，并没有置为NULL

```
154 struct inet_sock {
155     /* sk and pinet6 has to be the first two members of inet_sock */
156     struct sock      sk;
157     #if IS_ENABLED(CONFIG_IPV6)
158     struct ipv6_pinfo *pinet6;
159     #endif
160     /* Socket demultiplex comparisons on incoming packets. */
161     #define inet_daddr      sk->__sk_common.skc_daddr
162     #define inet_rcv_saddr  sk->__sk_common.skc_rcv_saddr
163     #define inet_dport      sk->__sk_common.skc_dport
164     #define inet_num        sk->__sk_common.skc_num
165
166     // ---- snip ----
167
168     be32      mc_addr;
169     struct ip_mc_socklist __rcu *mc_list;
170     struct inet_cork_full cork;
171 };
172
```

实例分析1

- 两个sock最终参与close()时，会发生double free

PoC

- `sockfd = socket(AF_INET, xx, IPPROTO_TCP);`
- `setsockopt(sockfd, SOL_IP, MCAST_JOIN_GROUP, xxxx, xxxx);`
- `bind(sockfd, xxxx, xxxx);`
- `listen(sockfd, xxxx);`
- `newsockfd = accept(sockfd, xxxx, xxxx); connect ()`
- `close(newsockfd) // first free (kfree_rcu)`
- `sleep(x) // wait rcu free(real free)`
- `close(sockfd) // double free`

Heap Spray 1

- uaf obj: **ip_mc_socklist** —> refill 点 —> kmalloc-64

```
struct ip_mc_socklist {  
    struct ip_mc_socklist __rcu *next_rcu;  
    struct ip_mreqn        multi;  
    unsigned int           sfmode;        /* MCAST_{INCLUDE,EXCLUDE} */  
    struct ip_sf_socklist __rcu *sflist;  
    struct rcu_head         rcu;  
};
```

Heap Spray 1

- rcu_head
 - 包含函数指针，利于PC指针劫持 (`rcu_process_callbacks`)


```
struct callback_head {  
    struct callback_head *next;  
    void (*func)(struct callback_head *head);  
};
```


Heap Spray 1

- refill obj 的选择

- 例如: ipv6_mc_socklist kmalloc-64 !!

用户态可控!



```
91 struct ipv6_mc_socklist {
92     struct in6_addr    addr;
93     int                ifindex;
94     struct ipv6_mc_socklist __rcu *next;
95     rwlock_t           sflock;
96     unsigned int        sfmode;        /* MCAST_{INCLUDE,
97     struct ip6_sf_socklist *sflist;
98     struct rcu_head      rcu;
99 };
```



Heap Spray 1

- ipv6_mc_list 之 addr 覆盖 ip_mc_socklist 之 next_rcu

```
struct ip_mc_socklist {  
    struct ip_mc_socklist __rcu *next_rcu;  
    struct ip_mreqn        multi;  
    unsigned int           sfmode;        /* MCAST_{INCLUDE,EXCLUDE} */  
    struct ip_sf_socklist __rcu *sflist;  
    struct rcu_head         rcu;  
};
```

Heap Spray 1

- spray 函数: `ipv6 setsockopt ()`
 - `kmalloc ipv6_mc_socklist`

实例分析2

- race导致的UAF和heap overflow: CVE-2017-7533

部分参考:Rooting Android 8 with a Kernel Space Mirroring Attack.pdf from Yong Wang

实例分析2

```
65 int inotify_handle_event(struct fsnotify_group *group,
66                         struct inode *inode,
67                         struct fsnotify_mark *inode_mark,
68                         struct fsnotify_mark *vfsmount_mark,
69                         u32 mask, void *data, int data_type,
70                         const unsigned char *file_name, u32 cookie)
71 {
72     struct inotify_inode_mark *i_mark;
73     struct inotify_event_info *event;
74     struct fsnotify_event *fsn_event;
75     int ret;
76     int len = 0;
77     int alloc_len = sizeof(struct inotify_event_info);
78
79     BUG_ON(vfsmount_mark);
80
81     if ((inode_mark->mask & FS_EXCL_UNLINK) &&
82         (data_type == FSNOTIFY_EVENT_PATH)) {
83         struct path *path = data;
84
85         if (d_unlinked(path->dentry))
86             return 0;
87     }
88     if (file_name) {
89         len = strlen(file_name);
90         alloc_len += len + 1;
91     }
92
93     // --- snip ---
94     event = kmalloc(alloc_len, GFP_KERNEL);
95     if (unlikely(!event))
96         return -ENOMEM;
97
98     fsn_event = &event->fse;
99     fsnotify_init_event(fsn_event, inode, mask);
100     event->wd = i_mark->wd;
101     event->sync_cookie = cookie;
102     event->name_len = len;
103     if (len)
104         strcpy(event->name, file_name);
```

strcpy(event->name, file_name)

file_name 可在其他线程用 reame
系统调用进行扩大名字长度从而导
致溢出!



实例分析2

- 但是有两种情况...
 - 短文件名触发 -> heap overflow
 - 长文件名触发 -> uaf

实例分析2

```
2418 static void copy_name(struct dentry *dentry, struct dentry *target)
2419 {
2420     struct external_name *old_name = NULL;
2421     if (unlikely(dname_external(dentry)))
2422         old_name = external_name(dentry);
2423     if (unlikely(dname_external(target))) {
2424         atomic_inc(&external_name(target)->u.count);
2425         dentry->d_name = target->d_name;
2426     } else {
2427         memcpy(dentry->d_iname, target->d_name.name,
2428             target->d_name.len + 1);
2429         dentry->d_name.name = dentry->d_iname;
2430         dentry->d_name.hash_len = target->d_name.hash_len;
2431     }
2432     if (old_name && likely(atomic_dec_and_test(&old_name->u.count)))
2433         kfree_rcu(old_name, u.head);
2434 }
```

长文件名走这里

短文件名走这里

old长文件名被kfree



Heap Spray 2

- 短文件名方式(目前很多机器不能用?)

- 故技新用：spray ipv6_mc_socklist

1. setsockopt() -> spray mc_list

2. 间隔close() -> free mc_list

3. open() -> fill hole



4. rename() -> overflow

kmalloc-64

Heap Spray 3

- 长文件名方式
 - 重点: 如何refill file_name 所在内存, 并最终扩大name长度呢
 - 大量线程进程race+spray

Heap Spray 3

- Spray 函数:长度和内容都可控, 此为佳
 - send , sendmsg
 - add_key (android已经不能用此方法了)

Heap Spray 3

- 探讨add_key

```
60 |SYSCALL_DEFINE5(add_key, const char __user *, _type,  
61 |    const char __user *, _description,  
62 |    const void __user *, _payload,  
63 |    size_t, plen,  
64 |    key_serial_t, ringid)  
65 |  
66 |    {  
67 |        key_ref_t keyring_ref, key_ref;  
68 |        char type[32], *description;  
69 |        void *payload;  
70 |        long ret;  
71 |        bool vm;  
72 |  
73 |        ret = -EINVAL;  
74 |        if (plen > 1024 * 1024 - 1)  
75 |            goto error;  
76 |  
77 |        /* draw all the data into kernel space */  
78 |        ret = key_get_type_from_user(type, _type, sizeof(type));  
79 |        if (ret < 0)  
80 |            goto error;  
81 |  
82 |        description = NULL;  
83 |        if (_description) {  
84 |            description = strndup_user(_description, KEY_MAX_DESC_SIZE);  
85 |            if (IS_ERR(description)) {  
86 |                ret = PTR_ERR(description);  
87 |                goto error;  
88 |            }  
89 |            if (!*description) {  
90 |                kfree(description);  
91 |                description = NULL;  
92 |            } else if ((description[0] == '.') &&  
93 |                (strncmp(type, "keyring", 7) == 0)) {  
94 |                ret = -EPERM;  
95 |            }  
96 |        }  
97 |  
98 |        if (ret < 0)  
99 |            goto error;  
100 |  
101 |        if (plen > 0) {  
102 |            if (vm) {  
103 |                ret = vm_fault_in_kernel(ringid, payload, plen);  
104 |                if (ret < 0)  
105 |                    goto error;  
106 |            }  
107 |            ret = keyctl_add_key(ringid, type, description, plen, 0);  
108 |            if (ret < 0)  
109 |                goto error;  
110 |        }  
111 |  
112 |        return 0;  
113 |    }  
114 |  
115 |error;  
116 |}
```

Heap Spray 3

- spray payload 1

```
101     vm = false;
102     if (payload) {
103         ret = -ENOMEM;
104         payload = kmalloc(plen, GFP_KERNEL | __GFP_NOWARN);
105         if (!payload) {
106             if (plen <= PAGE_SIZE)
107                 goto error2;
108             vm = true;
109             payload = vmalloc(plen);
110             if (!payload)
111                 goto error2;
112         }
113
114         ret = -EFAULT;
115         if (copy_from_user(payload, payload, plen) != 0)
116             goto error3;
117     }
118
```

Heap Spray 3

- spray payload 2 (18 bytes header+总量限制)

```
63 int user_preparse(struct key_prepared_payload *prep)
64 {
65     struct user_key_payload *upayload;
66     size_t datalen = prep->datalen;
67
68     if (datalen <= 0 || datalen > 32767 || !prep->data)
69         return -EINVAL;
70
71     upayload = kmalloc(sizeof(*upayload) + datalen, GFP_KERNEL);
72     if (!upayload)
73         return -ENOMEM;
74
75     /* attach the data */
76     prep->quotalen = datalen;
77     prep->payload[0] = upayload;
78     upayload->datalen = datalen;
79     memcpy(upayload->data, prep->data, datalen);
80     return 0;
81 }
82 EXPORT_SYMBOL_GPL(user_preparse);
```

Heap Spray 3

- spray description (较好)

```
278
279     if (desc) {
280         key->index key.desc len = descrlen;
281         key->index_key.description = kmemdup(desc, descrlen + 1, GFP_KERNEL);
282         if (!key->description)
283             goto no_memory_3;
284     }
```

Heap Spray 3

- 探讨其他的spray方法
 - msgsnd
 - 但是不利于CVE-2017-7533，可控数据前有可能有'\x00'截断
 - 也是比较好的spray函数

```

83 struct msg_msg *load_msg(const void __user *src, size_t len)
84 {
85     struct msg_msg *msg;
86     struct msg_msgseg *seg;
87     int err = -EFAULT;
88     size_t alen;
89
90     msg = alloc_msg(len);
91     if (msg == NULL)
92         return ERR_PTR(-ENOMEM);
93
94     alen = min(len, DATALEN_MSG);
95     if (copy_from_user(msg + 1, src, alen))
96         goto out_err;
97
98     for (seg = msg->next; seg != NULL; seg = seg->next) {
99         len -= alen;
100         src = (char __user *)src + alen;
101         alen = min(len, DATALEN_SEG);
102         if (copy_from_user(seg + 1, src, alen))
103             goto out_err;
104     }
105
106     err = security_msg_msg_alloc(msg);
107     if (err)
108         goto out_err;
109
110     return msg;
111
112 out_err:

```



Heap Spray 3

- 探讨其他的spray方法
 - setxattr等设置文件附加属性函数
 - 长度和内容可控
 - 可以通过单个文件或者大量文件进行spray
 - 对于CVE-2017-7533 race下的占位, 效果也比较好的

```

321 static long
322 setattr(struct dentry *d, const char __user *name, const void __user *value,
323         size_t size, int flags)
324 {
325     int error;
326     void *kvalue = NULL;
327     void *vvalue = NULL; /* If non-NULL, we used vmalloc() */
328     char kname[XATTR_NAME_MAX + 1];
329
330     if (flags & ~(XATTR_CREATE|XATTR_REPLACE))
331         return -EINVAL;
332
333     error = strncpy_from_user(kname, name, sizeof(kname));
334     if (error == 0 || error == sizeof(kname))
335         error = -ERANGE;
336     if (error < 0)
337         return error;
338
339     if (size) {
340         if (size > XATTR_SIZE_MAX)
341             return -E2BIG;
342         kvalue = kmalloc(size, GFP_KERNEL | __GFP_NOWARN);
343         if (!kvalue) {
344             vvalue = vmalloc(size);
345             if (!vvalue)
346                 return -ENOMEM;
347             kvalue = vvalue;
348         }
349         if (copy_from_user(kvalue, value, size)) {
350             error = -EFAULT;
351             goto out;
352         }
353         if ((strcmp(kname, XATTR_NAME_POSIX_ACL_ACCESS) == 0) ||
354             (strcmp(kname, XATTR_NAME_POSIX_ACL_DEFAULT) == 0))
355             posix_acl_fix_xattr_from_user(kvalue, size);
356     }

```



Heap Spray 3

- 探讨其他的spray方法
 - more !
 - 有优劣，视情况而定

Heap Spray 3

- race成功后
 - 采用spray+free iovec的方法产生内存holes，然后填充再溢出
 - 类似前面提到的

参考

- http://www.wowotech.net/memory_management/426.html
- Mirror Mirror: Rooting Android 8 with a Kernel Space Mirroring Attack (Yong Wang)
- https://github.com/hardenedlinux/offensive_poc
- <https://github.com/retme7/My-Slides>

End



PANGU TEAM